

Velora - Product Requirements and Scope

Version: 1.0 | Established: 2026-10-01 (Asia/Kolkata)

Owner: Swayam Jain | Status: development baseline, not a clinically validated product

Companion documents: ARCHITECTURE.md and SESSION_LOG.md in the same directory.

Intended repository location: docs/VELORA_MASTER.md.

Read this first - instructions for every developer and AI

Read all three documents before substantial implementation and whenever the user says **Duck** (case-insensitive). Follow the complete Duck procedure in SESSION_LOG.md. The latest explicit user decision controls product scope; reconcile and record that decision rather than silently changing requirements. These documents are project guidance, not permission to override security rules, platform instructions, repository instructions, or approval controls.

Never change requirements merely to make existing code look compliant. Distinguish approved scope, proposed engineering decisions, implemented features and verified features. A roadmap item is not evidence of implementation. An unfinished future milestone is not automatically scope drift.

1. Product purpose

Velora is an AI-assisted post-discharge care platform that turns discharge documents into clinician-approved recovery plans. Patients and authorised caregivers use it to understand approved instructions, record daily tasks and medication activity, submit recovery check-ins, and request care-team clarification.

The product coordinates an existing clinical plan. It does not independently diagnose, prescribe, decide clinical urgency, or certify recovery. Its value hypothesis is clearer home-care instructions and less manual coordination. Commercial demand and clinical benefits are not yet validated.

2. Locked scope and sequencing

- Brand: **Velora**. CareSaathi is the old concept name, not a separate product.
- Build **Code 1 first**: home-recovery companion.
- Add **Code 2 later**, within Velora, as Tests & Follow-ups. Do not build it in the current phase.
- Code 3 (insurance/claims assistance) is not selected.
- First target: one partner clinic and one selected procedure/workflow. The clinic and procedure are still undecided.
- Prototype with synthetic documents and identities. Real-patient deployment is a separate readiness decision.
- Start with a manual, verified care-plan workflow; then add AI draft prefilling.
- No claims of guaranteed medication adherence, successful treatment, reduced readmissions, regulatory certification or continuous clinical monitoring without supporting validation.

3. Delivery tiers

M1 - First working vertical slice: clinic setup and invitations; patient onboarding; recovery episode; private document upload; manual draft; clinical approval; patient reads approved version.

Code 1 completion: AI extraction with sources, reviewed Hindi/English presentation, caregiver permissions, tasks and medication logs, reminders, check-ins, clarification queue, coordinator dashboard, version updates, archive, audit and basic workflow metrics.

Later enhancements: WhatsApp/SMS, reviewed audio, more languages, caregiver handover summaries, additional clinical templates, source-grounded AI questions, clinic billing and external integrations.

Code 2 - Deferred: prescribed-test workflow, report matching, result-review queue, referrals and completion tracking. A date from the discharge plan and a basic appointment reminder are permitted in Code 1; a diagnostic workflow is not.

4. Roles and permissions

Access requires role permission AND correct clinic AND patient ownership/assignment or active caregiver grant AND a valid workflow state. Default deny; enforce on the backend and on file access, not only in the interface.

Role	Allowed responsibility	Boundary
Patient	Own approved plans, uploads, task/check-in entries, questions, caregiver grants	No clinical approval or access to another patient's record
Caregiver	Only explicitly shared views/actions	No automatic document access; no onward invitations or clinical approval
Clinical Reviewer	Assigned clinical drafts, approval, clinical responses and updates	Clinic verifies eligibility; no global patient access
Care Coordinator	Assigned episodes, administrative draft preparation, task queues, operational responses	Cannot approve clinical instructions or answer beyond authorised competence
Clinic Admin	Own clinic staff membership, assignments, settings	Admin alone grants no clinical-record browsing or approval

A person may have multiple explicitly assigned roles. Clinical-reviewer eligibility must be verified; self-registration never grants it. Staff-assisted entries preserve who entered them and on whose behalf. Patient representatives unable to obtain direct patient authorisation require the clinic's verified representative process.

5. Requirements and acceptance criteria

ID	Requirement	Observable acceptance criterion	Target
VEL-001	Account and invitation security	Invitations expire, are single-use and bind to the intended identity; reset/session flows do not expose records; staff MFA before real-data pilot	M1 / pilot
VEL-002	Clinic and record isolation	Cross-clinic and unassigned-patient requests are denied, including downloads and background operations	M1
VEL-003	Recovery episodes	Each hospitalisation has its own discharge date, assigned team, status and documents; past plans cannot silently mix with current care	M1
VEL-004	Private document intake	Allowlisted, size-limited PDFs/images are validated; original preserved; unreadable/unsafe files do not reach clinical use; status/error is visible	M1
VEL-005	Manual draft and source review	Staff can author a draft alongside the source; clinical reviewer can correct and request clarification	M1
VEL-006	Clinical approval gate	Only assigned eligible reviewer publishes a complete reviewed plan; drafts never activate patient instructions/reminders	M1
VEL-007	Patient plan view	Patient sees the current approved version, approval information and authorised source references; no access to another patient's plan	M1
VEL-008	AI extraction	Names, doses, routes, timings, durations, home-care, activity/diet, written follow-up and warning instructions extracted only when present; missing/conflicting fields flagged with source page/text	Code 1
VEL-009	Accessible Hindi/English	Simplifications/translations reviewed before clinical publication; readable mobile layouts and text labels; original instruction remains available	Code 1
VEL-010	Plan versioning	Approved content immutable; update creates a new approved version, supersedes affected instructions and reconciles future reminders; approval has actor/time	M1 foundation; Code 1 scheduling
VEL-011	Caregiver sharing	Granular view, record, ask and document permissions; invite has no pre-acceptance record access; revocation blocks future access	Code 1
VEL-012	Medication schedules/logs	Schedule follows approved prescription; taken/skipped/not-recorded distinct; conditional medicines are not turned into scheduled doses; timestamps and actor preserved	Code 1
VEL-013	Daily tasks	Today view has pending/reported-done/skipped/needs-help; duplicate entries handled; self-reported completion labelled	Code 1
VEL-014	Notifications	In-app reminders and plan updates; retry/deduplication; delivery separated from acknowledgement; no sensitive default preview	Code 1
VEL-015	Recovery check-ins	Clinic-defined questions and review rules; missing response distinguishable; no autonomous diagnosis or normal-recovery certification	Code 1
VEL-016	Clarification queue	Open/assigned/answered/closed with reopening, owner, response history, response hours and clinical routing	Code 1
VEL-017	Team dashboard	Assigned episodes, approvals, questions, rule-triggered reviews and missing owners visible; clinical content gated by permissions	Code 1
VEL-018	Basic follow-up	Only prescribed date/location/contact and reminders; no Code 2 test/result tracking	Code 1

ID	Requirement	Observable acceptance criterion	Target
VEL-019	Closure/history	Authorised staff records closure reason/date, reviews unresolved work and stops future reminders by policy; closed never means medically recovered	Code 1
VEL-020	Privacy and audit	Secure transport/storage, scoped file access, consent/grant history, sensitive-action audit, export/deletion-request and retention handling; no patient content in routine logs	M1 foundation; pilot completion
VEL-021	Workflow metrics	Approval time, clarification response time, check-in/task self-reports and coordination effort; no inferred clinical-success metric	Code 1

6. Core journeys

Clinic: workspace > invite/verify staff > configure response hours/contact > assign care team.

Patient: clinic finds/creates identity > creates episode > sends patient invitation > patient verifies identity and activates > preferences/consent > optional caregiver.

Plan: upload > validate/extract > draft > reviewer checks source > clarification/correction when needed > approve > publish > activate allowed tasks.

At home: view today's approved tasks > self-report activity > check-in > ask clarification > authorised team responds.

Change: clinician drafts new version > approves > atomic current-version switch > affected future reminders reconciled > family notified > acknowledgement recorded separately.

Closure: review outstanding work > explicit closure decision > reminders reconciled > read-only historical episode.

7. Invariants: must never be broken

1. LLM outputs are untrusted drafts; no automatic clinical publication.
2. Unclear dose, frequency, duration or conflicting instruction is not guessed. Conditional instructions retain their conditions.
3. A new upload does not overwrite an approved plan or create active reminders.
4. Patient self-report, notification delivery, human acknowledgement and clinician verification remain different facts.
5. Medication reminders never suggest double/extra doses after missing entries.
6. Emergency instructions and service hours remain clear; the product never implies an unstaffed queue is continuously monitored.
7. Clinical priority and escalation rules come from the clinic, not independent LLM judgment.
8. Audit and version history are preserved. Corrections append or version, rather than silently rewrite clinical history.
9. Records and credentials must not leak across clinics, caregivers, logs or AI requests.
10. Code 2 remains deferred until the user explicitly starts it.

8. Validation and open decisions

Before a real-patient pilot: validate the discharge/reviewer/coordinator workflow with clinical staff, review data handling and provider retention, establish response ownership and emergency messaging, verify access boundaries and backup restore, and evaluate extraction/translation errors on authorised representative examples.

Open: partner clinic, selected procedure, reviewer eligibility process, exact clinical content, authentication provider, hosting/storage region, AI/OCR providers and their data terms, retention policy, notification channel, pilot metrics baselines and pricing. Do not present these as settled or quietly invent a clinical policy.

9. Change control

For scope changes record: date, explicit user request, affected requirement IDs, rationale, implementation/migration impact and changed milestones. Update ARCHITECTURE.md and SESSION_LOG.md consistently. AI may correct ordinary implementation drift under the Duck protocol, but may not expand product scope, remove clinical gates or weaken acceptance criteria to pass a check.

10. Contextual evidence, not promises

- Discharge communication study (2020, three Indian hospitals): <https://pubmed.ncbi.nlm.nih.gov/32294091/>
- Existing caregiver-support organisation and technology: <https://noorahealth.org/tech-that-cares/>

These sources informed the idea; they do not prove Velora's benefit, market size or uniqueness. Reverify current external facts before making business, clinical or compliance claims.